

Artifact 8 — Prioritise Assumptions

Risk & Ethics Gating · Stage 4

DOCUMENT DETAILS

Project	Home-Care-AI	Stage	Discovery → Stage 4 (Risk & Ethics Gating)
Date	2026-03-26	Beachhead	Care Coordinator (CC)
Methodology	ICE scoring (Impact × Confidence × Ease) → Impact × Risk matrix → Experiment queue		
Input	Artifact 7 — 45 assumptions across 8 categories		
Feeds into	Artifact 9 — Ethics Trust Mapping (Skill 9); Artifact 11 — Brainstorm Experiments (Skill 11)		
Amended	2026-03-26 — Fix 1: removed duplicate V-2 row; Fix 2: corrected GTM-4 ICE math in Section 5 (Confidence 9→8, score 648→576); Fix 3: removed E-3 from Q2 table (Confidence=8 places it in Q4 only)		

Table of Contents

Table of Contents	2
1. Scoring Framework	3
ICE Formula	3
Impact × Risk Matrix	3
2. Full ICE Scoring — All 45 Assumptions	3
3. Impact × Risk Matrix — Summary	6
Q1 — High Impact + High Risk → DESIGN EXPERIMENT NOW	6
Q2 — High Impact + Medium Risk → DESIGN EXPERIMENT (lower urgency)	6
Q3 — Medium Impact → MONITOR DURING E1/E2/E3 OR DEFER	7
Q4 — High Confidence → PROCEED (no experiment needed)	8
4. Priority Experiment Queue	8
Tier 1 — Before Any Code Is Written (Immediate)	8
Tier 2 — Week 1–2 (Launch E1 + E2 + E3 Concurrently)	9
Tier 3 — Alongside E1/E2/E3 (Parallel Research)	9
5. ICE Math — Key Assumptions (Top 20, Shown in Full)	10
6. Critical Path Summary	10
7. Assumptions Explicitly Deferred	11
8. Downstream Feed	11
9. Handshake Table	12
10. Action Items	12

1. Scoring Framework

1. Scoring Framework

ICE Formula

ICE = Impact × Confidence × Ease

All dimensions rated 1–10. Raw ICE score range: 1–1,000.

Dimension	Definition	Scale
Impact	Consequence of this assumption being wrong — how much does it affect the product's ability to deliver value, reach product-market fit, or avoid legal/ethical harm?	10 = product cannot ship / collapses; 1 = negligible
Confidence	Current confidence that this assumption is correct, based on evidence from Artifacts 2a–2d and 3	Low = 2–3; Medium = 5–6; High = 8–9
Ease	How quickly and cheaply can this assumption be tested?	10 = 1 day, near-zero cost; 1 = requires 6+ months or prohibitive cost

Impact × Risk Matrix

Risk is the inverse of Confidence: High Risk = Low Confidence (unknown = risky).

Quadrant	Impact	Risk	Action
Q1 — Design Experiment	High (8–10)	High (Low confidence, ≤ 4)	Build an experiment immediately — critical unknowns blocking progress
Q2 — Experiment (lower urgency)	High (8–10)	Medium (Confidence 5–6)	Design an experiment — important, plausible, but not at maximum urgency
Q3 — Monitor or Research	Medium (6–7)	Any	Research, observe during E1/E2/E3, or defer until Q1/Q2 resolved
Q4 — Proceed	Any	Low (High confidence, ≥ 7)	No experiment needed — treat as confirmed; implement the mitigation

2. Full ICE Scoring — All 45 Assumptions

2. Full ICE Scoring — All 45 Assumptions

ID	Category	Assumption (Short)	Impact	Confidence	Ease	ICE	Quadrant
E-1	Ethics	SPP matching does not create discriminatory carer selection	10	3	4	120	Q1 ■■ CRITICAL
E-3	Ethics	Family notification never precedes patient notification	10	8	5	400	Q4 ■■ CRITICAL — Structural gate
V-2	Value	Coordinator trusts machine ranking without calling to verify	9	5	7	315	Q2

ID	Category	Assumption (Short)	Impact	Confidence	Ease	ICE	Quadrant
U-1	Usability	3-tap flow completable at 6:30 AM without training	9	5	8	360	Q2
VI-4	Viability	APP 8 resolved by design + privacy notice (not per-client consent)	9	5	5	225	Q2 ■■■ CRITICAL
V-1	Value	SPP match produces better clinical/relational outcomes vs. availability-only	9	5	5	225	Q2
F-1	Feasibility	Staff update availability reliably in app (vs. calling coordinator)	9	3	7	189	Q1
VI-1	Viability	Agencies pay for standalone tool without EMR integration at launch	9	3	5	135	Q1 ■■■ CRITICAL
S-1	Strategy	AlayaCare does not add soft preference matching within 12–18 months	9	5	7	315	Q2
GTM-1	Go-to-Market	CC is both champion and veto player for agency purchase	8	8	8	512	Q4
GTM-4	Go-to-Market	30-day trial includes enough incidents for genuine evaluation	8	8	9	576	Q4
T-4	Team	Angela + Tom engagement sustainable for 4+ months	8	5	8	320	Q2
GTM-3	Go-to-Market	E1 participants become willing reference customers	8	5	7	280	Q2
F-2	Feasibility	Rule-based SPP score correlates with coordinator judgment $r > 0.70$	8	5	6	240	Q2
E-5	Ethics	SPP fields classified correctly under APP (personal/sensitive information)	8	5	5	200	Q2
S-2	Strategy	CC beachhead unlocks HCN/FC without separate sales motion	8	5	5	200	Q2
VI-2	Viability	Per-agency subscription covers COGS at target margin	8	3	6	144	Q1
VI-3	Viability	Australian home care market large enough for 24-month growth	8	3	6	144	Q1
E-2	Ethics	Matching algorithm does not amplify coordinator bias in SPP data	8	3	6	144	Q1
S-3	Strategy	SPP data moat is deep enough to prevent switching after 6 months	8	3	6	144	Q1
V-5	Value	Agency owners perceive coordinator departure as material business risk	8	3	5	120	Q1
GTM-6	Go-to-Market	Concierge-to-product transition does not erode coordinator trust	8	3	5	120	Q1
T-2	Team	Build Lambda + SPP + WhatsApp to prototype in 8 weeks	8	3	7	168	Q1
VI-5	Viability	SPP moat drives monthly churn $< 2\%$ post-3 months	8	3	2	48	Q1 — Defer (not testable pre-launch)

ID	Category	Assumption (Short)	Impact	Confidence	Ease	ICE	Quadrant
E-4	Ethics	Carer availability/location data ≠ GPS surveillance under APP	7	5	5	175	Q3
V-3	Value	SPP has standalone value before matching engine is built	7	5	7	245	Q3
U-2	Usability	3 candidates is the right decision surface (not 1, not 5)	7	5	8	280	Q3
U-3	Usability	Familiarity flag is the trust-producing field on the candidate card	7	5	8	280	Q3
F-5	Feasibility	WhatsApp reply parsing reliable for single-character approval commands	7	5	7	245	Q3
T-1	Team	Domain knowledge sufficient without dedicated clinical advisor	7	5	7	245	Q3
GTM-2	Go-to-Market	Peer referral is primary acquisition channel for agency owners	7	5	7	245	Q3
S-5	Strategy	Product positioned as clinical intelligence layer, not scheduling tool	7	8	9	504	Q4
T-3	Team	Privacy counsel available on retainer basis (no full-time hire required)	7	8	9	504	Q4
V-4	Value	Carer briefing notification reduces post-assignment phone calls	7	3	7	147	Q3
U-5	Usability	Onboarding wizard produces accurate SPP without high correction rate	7	3	6	126	Q3
V-6	Value	Family notification prevents client churn at agency level	7	5	6	210	Q3
U-4	Usability	Tom and Angela served by same interface without product variants	6	5	7	210	Q3
VI-6	Viability	NDIS compliance doesn't create adoption-blocking burden	6	3	6	108	Q3
F-3	Feasibility	Google Maps API accurate for regional Australian matching	6	8	9	432	Q4
F-4	Feasibility	DynamoDB SPP query < 200ms at 60-client scale	6	8	9	432	Q4
F-6	Feasibility	Twilio SMS reliable for international family contacts	5	7	8	280	Q4
S-4	Strategy	Regulatory environment stable for 24-month build horizon	6	5	5	150	Q3
E-6	Ethics	SPP does not create "performing wellness" dynamic for clients	6	5	3	90	Q3 — Defer
GTM-5	Go-to-Market	"Bus-proof" messaging resonates more than efficiency messaging	6	3	5	90	Q3 — Defer
T-5	Team	Team applies 75% failure rate honestly to experiment results	7	3	4	84	Q3 — Defer

3. Impact × Risk Matrix — Summary

3. Impact × Risk Matrix — Summary

Q1 — High Impact + High Risk → DESIGN EXPERIMENT NOW

Q1 — High Impact + High Risk → DESIGN EXPERIMENT NOW

These are the most dangerous unknowns. High consequence of being wrong; low current confidence.

Priorit y	ID	Assumption	ICE	Experiment
1	E-1 ■■■	SPP matching ≠ discriminatory carer selection	120	Legal opinion — anti-discrimination law (Sex Discrimination Act, Disability Discrimination Act) + employment law review
2	VI-1 ■■■	Agencies pay without EMR integration at launch	135	Interview 2 agency owners — past behaviour: "Tell me about the last tool you adopted without integration"
3	F-1	Staff update availability reliably in app	189	E1 observation — track how many absences are self-reported vs. coordinator-discovered
4	T-2	Build prototype in 8 weeks	168	Sprint 0 sizing — engineering 3-point estimate for ENG-1 + ENG-3
5	VI-2	Per-agency subscription covers COGS	144	Unit economics model — AWS Lambda + DynamoDB + Twilio cost at 3 incidents/week/agency
6	VI-3	Australian home care market large enough	144	Secondary research — ACFA report, NDIS provider register, My Aged Care data
7	E-2	Matching doesn't amplify coordinator bias	144	E2 probe — for each SPP entry, ask: "Did Mrs. Kim tell you this, or is this your inference?"
8	S-3	SPP moat deep enough to prevent switching	144	E2 rebuild probe — "If we deleted the SPP tomorrow, how long to rebuild?"
9	V-5	Agency owners see coordinator departure as material risk	120	Interview 2 agency owners — Mom Test: "When did you last lose a coordinator? What happened?"
10	GTM-6	Concierge-to-product transition doesn't erode trust	120	Blind transition test at E1 week 3 — replace human team with algorithm, observe acceptance rate
—	VI-5	SPP moat drives churn < 2%	48	■■■ DEFERRED — not testable before 6 months of live data. Design 6-month cohort metric at launch

Q2 — High Impact + Medium Risk → DESIGN EXPERIMENT (lower urgency)

Q2 — High Impact + Medium Risk → DESIGN EXPERIMENT (lower urgency)

Important assumptions with plausible evidence — test after Q1 experiments are running.

Priorit y	ID	Assumption	ICE	Experiment
1	GTM-4	30-day trial has enough incidents	576	Confirmed by CC interviews. Track incident rate during E1. No separate experiment needed.
2	GTM-1	CC is champion and veto player	512	Confirmed by CC interviews. Observe in E1. No separate experiment needed.
3	U-1	3-tap flow completable at 6:30 AM without training	360	E3 (Figma prototype) — unguided walkthrough, 5 scenarios, measure time-to-completion
4	T-4	Angela + Tom engagement sustainable 4+ months	320	Obtain written consent from coordinator + agency management before E1 begins

Priorit y	ID	Assumption	ICE	Experiment
5	V-2	Coordinator trusts machine ranking to approve	315	E3 (Figma prototype) — observe acceptance rate of top-ranked candidate
6	S-1	AlayaCare doesn't add soft preference in 12–18 months	315	Competitive intelligence — monitor AlayaCare product updates every 6 weeks
7	GTM-3	E1 participants become reference customers	280	Ask at E1 end: "Would you speak with other coordinators about your experience?"
8	F-2	Rule-based SPP score correlates with coordinator judgment $r > 0.70$	240	E1 — compare algorithm ranking to coordinator's manual choice for each incident
9	VI-4 ■■	APP 8 resolved by design + privacy notice	225	Privacy counsel review — 2-hour retainer: classify data flows, confirm notice-vs-consent path
10	V-1	SPP match produces better outcomes than availability-only	225	E1 — track distress events, refusals, complaints during SPP-matched vs. unmatched periods
11	E-5	SPP fields classified correctly under APP	200	Privacy counsel review — classify each SPP field as personal/sensitive/operational under APP 3
12	S-2	CC beachhead unlocks HCN/FC without separate sales motion	200	Ask Angela + Tom at 3-month E1 review: "Would you want nurses to use this? How would you sell it?"

Q3 — Medium Impact → MONITOR DURING E1/E2/E3 OR DEFER

Q3 — Medium Impact → MONITOR DURING E1/E2/E3 OR DEFER

Worthwhile to track; not worth a dedicated experiment at this stage.

ID	Assumption	ICE	Action
V-3	SPP standalone value before matching engine	245	Probe during E2 — WTP question at session end
U-2	3 candidates right decision surface	280	E3 variant test — run 3 vs. 5 candidates in 2 scenario sets
U-3	Familiarity flag is trust-producing field	280	E3 think-aloud — "What do you look at first?"
F-5	WhatsApp reply parsing reliable	245	Technical spike — 50 synthetic reply variants
T-1	Domain knowledge sufficient without clinical advisor	245	E2 probe — ask Angela: "Are there SPP fields this template misses?"
GTM-2	Peer referral primary acquisition channel	245	Ask in E1 — "How did you find the last tool you adopted?"
E-4	Carer location data ≠ GPS surveillance	175	Legal/HR review — brief opinion on Fair Work + Privacy Act
V-4	Briefing notification reduces phone calls	147	E1 observation — does coordinator call carer after notification sent?
U-5	Onboarding wizard NLP accuracy	126	E2 comparison — paper form vs. NLP extraction on same clients
V-6	Family notification prevents churn	210	E1 follow-up — track retention of notified families over 4 weeks
U-4	Same interface for Tom and Angela	210	E1 observation — note any workflow differences between coordinators
VI-6	NDIS compliance doesn't block adoption	108	1 interview — ask 1 NDIS-registered agency manager about software compliance checks
S-4	Regulatory environment stable 24 months	150	Regulatory briefing — 1-hour with aged care compliance consultant

ID	Assumption	ICE	Action
E-6	No "performing wellness" dynamic for clients	90	■■ DEFERRED — longitudinal indicator only; check at 6-month review
GTM-5	"Bus-proof" messaging resonates more than efficiency	90	■■ DEFERRED — landing page A/B test when product is closer to launch
T-5	Team applies 75% failure rate honestly	84	■■ DEFERRED — observable only after experiment results come in

Q4 — High Confidence → PROCEED (no experiment needed)

Q4 — High Confidence → PROCEED (no experiment needed)

Confirmed by evidence or are structural requirements. Implement the mitigation directly.

ID	Assumption	ICE	Required Action
GTM-4	30-day trial has enough incidents	576	Track incident rate during E1 — confirmed by CC interview frequency data
GTM-1	CC is champion and veto player	512	Design product and sales motion around coordinator as champion — confirmed by CC personas
S-5	Clinical intelligence framing, not scheduling tool	504	Internal alignment check every sprint — confirmed by Artifact 3 strategic implication
T-3	Privacy counsel on retainer basis	504	Obtain 3 quotes from Australian APP health data privacy firms this week
F-3	Google Maps API accurate for regional AU	432	Technical spike — 10 address pairs — confirmed expectation; spike is a formality
F-4	DynamoDB SPP query < 200ms at 60-client scale	432	Technical spike — confirm; trivial at this scale
E-3	Family notification never precedes patient	400	■■ Structural gate — not an experiment. Must appear in agentic-logic-spec pseudocode as an explicit coordinator_approved = true prerequisite gate. Cannot be overridden by any user request.
F-6	Twilio SMS reliable for international contacts	280	Twilio documentation review + pricing model — confirm UK/NZ delivery and rates

4. Priority Experiment Queue

4. Priority Experiment Queue

Ranked by ICE score within each urgency tier. This is the sequence in which experiments should be designed and launched.

Tier 1 — Before Any Code Is Written (Immediate)

Tier 1 — Before Any Code Is Written (Immediate)

#	ID	Assumption to Test	Method	Timeline	Success Threshold
1	E-1	SPP matching ≠ discriminatory carer assignment	Legal opinion — Anti-Discrimination Act + employment law	1 week	Written legal opinion: matching on client preference is lawful personal care exception OR specifies mitigation required

#	ID	Assumption to Test	Method	Timeline	Success Threshold
2	VI-4 + E-5	APP 8 + SPP data classification	Privacy counsel retainer review (2 hrs)	1 week	(a) Privacy notice update sufficient for APP 8; (b) each SPP field classified as personal/sensitive/operational
3	T-3	Privacy counsel available on retainer	Obtain 3 quotes from AU privacy law firms	3 days	≥ 1 firm confirms availability + cost < \$5,000 AUD for initial review
4	T-4	Angela + Tom 4-month engagement	Written consent from coordinator + agency management	Before E1 start	Signed consent from both coordinators and their agency owners
5	T-2	8-week build timeline achievable	Sprint 0 — engineering 3-point estimate (ENG-1 + ENG-3)	3 days	Likely estimate ≤ 6 weeks for ENG-1 + ENG-3; if > 6 weeks, revise scope

Tier 2 — Week 1–2 (Launch E1 + E2 + E3 Concurrently)

Tier 2 — Week 1–2 (Launch E1 + E2 + E3 Concurrently)

#	ID	Assumption to Test	Method	Timeline	Success Threshold
6	F-1	Staff update availability reliably	E1 — track absence self-report rate vs. coordinator-discovered	2 weeks	≥ 70% of absences reported by carer proactively; < 30% discovered by coordinator
7	V-1 + V-2	SPP match quality + coordinator trust	E1 — track distress events; E3 — top-candidate acceptance rate	2 weeks / 1 day	E1: 0 Henderson-equivalent events in matched cohort; E3: ≥ 70% top-candidate acceptance
8	V-5 + VI-1	Agency owner bus risk + standalone WTP	Interview 2 agency owners (Mom Test)	1 week	≥ 1 owner confirms coordinator departure is a "significant risk" unprompted; ≥ 1 confirms willingness to trial without EMR integration
9	E-2	SPP bias source (direct vs. inferred)	E2 — for each SPP entry, ask: "Did [client] tell you this directly?"	During E2	≥ 80% of SPP entries are direct client statements, not coordinator inference
10	F-2	Rule-based score correlates with judgment	E1 — show algorithm ranking to coordinator post-match; measure divergence	2 weeks	≥ 70% algorithm top pick matches coordinator's manual choice

Tier 3 — Alongside E1/E2/E3 (Parallel Research)

Tier 3 — Alongside E1/E2/E3 (Parallel Research)

#	ID	Assumption to Test	Method	Timeline	Success Threshold
11	VI-2 + VI-3	Unit economics + market size	Cost model build + secondary research (ACFA, NDIS data)	1 week	Gross margin > 60% at target subscription price; ≥ 500 agencies in AU at 20–200 client scale
12	S-3	SPP moat depth	E2 rebuild probe: "How long to rebuild this knowledge without the system?"	During E2	Answer > 4 weeks = deep moat; < 2 weeks = shallow moat requiring revised strategy
13	GTM-6	Concierge-to-product trust transition	E1 week 3 blind transition — replace human team with algorithm	Week 3	Coordinator acceptance rate does not drop by > 10 percentage points post-transition

#	ID	Assumption to Test	Method	Timeline	Success Threshold
14	S-1	AlayaCare competitive window	Monitor AlayaCare release notes, conference announcements	Ongoing (every 6 weeks)	No soft preference feature announcement within 12 months = window confirmed
15	GTM-3	E1 participants as reference customers	Ask at E1 close: "Would you speak with other coordinators about your experience?"	End of E1	≥ 1 coordinator confirms willingness to be a named reference

5. ICE Math — Key Assumptions (Top 20, Shown in Full)

5. ICE Math — Key Assumptions (Top 20, Shown in Full)

Per CLAUDE.md Article III: "show ICE math."

ID	Impact	×	Confidence	×	Ease	=	ICE	Quadrant
GTM-4	8	×	8	×	9	=	576	Q4 — Proceed
GTM-1	8	×	8	×	8	=	512	Q4 — Proceed
S-5	7	×	8	×	9	=	504	Q4 — Proceed
T-3	7	×	8	×	9	=	504	Q4 — Proceed
F-3	6	×	8	×	9	=	432	Q4 — Proceed
F-4	6	×	8	×	9	=	432	Q4 — Proceed
E-3	10	×	8	×	5	=	400	Q4 — Structural gate
U-1	9	×	5	×	8	=	360	Q2 — Experiment
T-4	8	×	5	×	8	=	320	Q2 — Experiment
V-2	9	×	5	×	7	=	315	Q2 — Experiment
S-1	9	×	5	×	7	=	315	Q2 — Monitor
GTM-3	8	×	5	×	7	=	280	Q2 — Experiment
F-2	8	×	5	×	6	=	240	Q2 — Experiment
V-3	7	×	5	×	7	=	245	Q3 — Probe
VI-4	9	×	5	×	5	=	225	Q2 — Legal review
V-1	9	×	5	×	5	=	225	Q2 — Experiment
F-1	9	×	3	×	7	=	189	Q1 — Experiment
T-2	8	×	3	×	7	=	168	Q1 — Sprint 0
VI-2	8	×	3	×	6	=	144	Q1 — Model
E-1	10	×	3	×	4	=	120	Q1 — Legal opinion

Note: E-1 scores 120 (low ICE) despite being CRITICAL because Ease = 4 (legal opinion takes time and cost). Low ICE ≠ low urgency for Critical items — the matrix quadrant (Q1: High Impact + High Risk) governs priority, not ICE score alone.

6. Critical Path Summary

6. Critical Path Summary

The 4 Critical assumptions must resolve before any of these components enter build

Critical Assumption	Blocks	Resolution Owner	Deadline
E-1 — SPP ≠ discriminatory selection	Matching engine (ENG-1 + S1 + S2 + S3)	Legal counsel	Before Sprint 1
E-3 — Family notification ≠ patient-first violation	Notification pipeline (S6 + ENG-4)	Engineering — agentic-logic-spec	Before Sprint 1
VI-4 — APP 8 resolved by design	WhatsApp integration (ENG-2)	Privacy counsel + PM Lead	Before Sprint 1
VI-1 — Agencies pay without EMR integration	Go-to-market / sales motion	PM Lead — 2 agency owner interviews	Before launch planning

If E-1 returns a negative legal opinion (i.e., automated SPP-based filtering is discriminatory), the matching engine must be redesigned: SPP becomes advisory (shown to coordinator but not used in ranking), and all filtering remains under coordinator judgment. This does not kill the product — it changes the automation level from auto-rank to auto-present.

If VI-1 returns a negative (agencies require EMR integration before paying), the go-to-market strategy must shift: either build an AlayaCare integration before launch (6–9 months delay) or target the subset of agencies not currently on any platform (greenfield agencies).

7. Assumptions Explicitly Deferred

7. Assumptions Explicitly Deferred

These have been scored and categorised. They will not receive experiments at this stage.

ID	Assumption	ICE	Reason for Deferral
VI-5	SPP moat drives churn < 2%	48	Not testable before 6 months of live product data — design cohort metric at launch
E-6	No "performing wellness" dynamic	90	Longitudinal signal only — observable at 6-month review, not pre-launch
GTM-5	"Bus-proof" messaging resonates more	90	Landing page A/B test appropriate when product is 4–6 weeks from launch
T-5	Team applies 75% failure rate honestly	84	Observable only after first experiment results arrive — revisit post-E1

8. Downstream Feed

8. Downstream Feed

This prioritisation directly determines what Skill 9 (ethics-trust-mapping) and Skill 11 (brainstorm-experiments-new) receive:

→ ethics-trust-mapping (Skill 9)

- E-1: SPP data fields that enable carer filtering → Green/Yellow/Red classification
- E-2: SPP data bias risk → audit protocol for preference data capture
- E-3: Notification pipeline ordering → structural gate requirement
- E-4: Carer location data → data sensitivity classification
- E-5: All SPP fields → APP sensitivity classification
- VI-4: WhatsApp message content → data minimisation check

→ agentic-safety-discovery (Skill 10)

- F-1: Staff availability update → Level 1/2/3 classification for availability-stale event
- F-2: SPP match confidence scoring → confidence threshold constants
- E-3: Notification gate → Level 2/3 HITL requirement before family notification

→ brainstorm-experiments-new (Skill 11)

Top 15 experiments from Priority Experiment Queue (Sections 4 Tier 1–3) form the direct input. E1, E2, E3 designs from Artifact 5 are confirmed as the primary prototypes.

9. Handshake Table

9. Handshake Table

ID	From Artifact	To Next Skill	What Transfers
HS-DISC-PR I-01	Artifact 8 — Q1 Critical assumptions (E-1, VI-1, VI-4)	ethics-trust-mapping (Skill 9)	Critical ethics and viability risks → data classification + consent design
HS-DISC-PR I-02	Artifact 8 — E-3 structural gate	agentic-logic-spec (Execution Plugin)	coordinator_approved = true required before family_notification_triggered = true
HS-DISC-PR I-03	Artifact 8 — Priority Experiment Queue Tier 1–3	brainstorm-experiments-new (Skill 11)	Top 15 experiments → XYZ hypothesis + prototype design
HS-DISC-PR I-04	Artifact 8 — Q4 confirmed assumptions	startup-canvas (Strategy Plugin)	Confirmed strategic facts (GTM-1 CC champion, S-5 clinical intelligence framing) → canvas inputs
HS-DISC-PR I-05	Artifact 8 — VI-2 unit economics gap	ai-unit-economics (Strategy Plugin)	Cost model required — feeds into subscription pricing + gross margin projection

10. Action Items

10. Action Items

Date	Owner	Action
2026-03-27	PM Lead	Engage privacy counsel — APP 8 + anti-discrimination law review (E-1, VI-4, E-4, E-5) — obtain 3 quotes (T-3)
2026-03-27	PM Lead	Run ethics-trust-mapping (Skill 9) — feed E-1, E-2, E-3, E-4, E-5 + VI-4 into Green/Yellow/Red classification
2026-03-27	PM Lead	Secondary market research — Australian home care agency count at 20–200 client scale (VI-3)
2026-03-28	PM Lead	Interview 2 agency owners using Mom Test (V-5 + VI-1) — recruit via Angela and Tom referral
2026-03-28	Engineer	Sprint 0 sizing — ENG-1 + ENG-3 3-point estimate (T-2)
2026-03-28	Engineer	Technical spikes — Google Maps API (F-3) + DynamoDB benchmark (F-4)
2026-03-28	PM Lead	Unit economics model — AWS Lambda + DynamoDB + Twilio at 3 incidents/week/agency (VI-2)
2026-03-29	PM Lead	Obtain written consent from Angela + Tom + their agency management (T-4)
2026-03-29	Designer	Begin E3 Figma prototype — 3-tap approval card, 5 vacancy scenarios (U-1 + V-2)
2026-04-03	PM Lead	Launch E1 (concierge prototype) — 2-week run with Angela + Tom

Date	Owner	Action
2026-04-03	PM Lead	Run E2 (preference extraction session) with Angela — 90 min, 10 clients
2026-04-07	PM Lead	Run agentic-safety-discovery (Skill 10) after ethics-trust-mapping completes
2026-04-07	PM Lead	Run brainstorm-experiments-new (Skill 11) using Priority Experiment Queue Tier 1–3

Home-Care-AI · Artifact 8 · Prioritise Assumptions · 2026-03-26