Paweł Huryn

AI PM Learning Roadmap

1 WTF is AI PM?

2 Foundations

- Neural networks, transformers
- LLMs and other architectures

3 Prompt Engineering

- Prompting principles
- High-ROI use cases

4 Context Engineering

- Context windows
- Types of context
- Information flow inside the context
- Retrieving, compressing, and saving context

5 RAG for PMs

- Types of RAG
- Vector stores, embeddings
- Retrieval & ranking
- RAG architectures, limitations

6 Model Interfaces & API

- Assistants & Responses API
- Requests, responses, tokens
- Latency, cost, reliability tradeoffs
- Interface constraints PMs must account for

7 How LLMs Learn & Adapt

- Supervised fine-tuning
- Preference optimization (DPO, ORPO, KTO)
- Reinforcement learning (RL) intuition
- How to choose a fine-tuning approach

8 AI Evals & Observability

- AI observability & eval platforms
- AI evals: error analysis process
- Failure modes, types, automatic evals
- Why, when, and how to measure human-model agreement, TPR

9 AI Agents for PMs

- LLM workflows vs. agents vs. multi-agent systems & architectures
- Planning, reflection, adaptation
- Tools, memory types, MCP
- Guardrails & evals for AI agents

10 AI Product Strategy

The Product Compass Newsletter